


# Journal Pre-proof

Learning dynamics from large biological datasets: machine learning meets systems biology

William Gilpin, Yitong Huang, Daniel B. Forger



PII: S2452-3100(20)30014-7

DOI: <https://doi.org/10.1016/j.coisb.2020.07.009>

Reference: COISB 286

To appear in: *Current Opinion in Systems Biology*

Received Date: 4 June 2020

Revised Date: 2 July 2020

Accepted Date: 20 July 2020

Please cite this article as: Gilpin W, Huang Y, Forger DB, Learning dynamics from large biological datasets: machine learning meets systems biology, *Current Opinion in Systems Biology*, <https://doi.org/10.1016/j.coisb.2020.07.009>.

This is a PDF file of an article that has undergone enhancements after acceptance, such as the addition of a cover page and metadata, and formatting for readability, but it is not yet the definitive version of record. This version will undergo additional copyediting, typesetting and review before it is published in its final form, but we are providing this version to give early visibility of the article. Please note that, during the production process, errors may be discovered which could affect the content, and all legal disclaimers that apply to the journal pertain.

© 2020 Published by Elsevier Ltd.

**Title: Learning dynamics from large biological datasets: machine learning meets systems biology**William Gilpin<sup>1</sup>, Yitong Huang<sup>2</sup>, Daniel B Forger<sup>1,3,4,5</sup>

1. Quantitative Biology Initiative, Harvard University
2. Department of Mathematics, Dartmouth College
3. Department of Mathematics, University of Michigan
4. Department of Computational Medicine and Bioinformatics, University of Michigan
5. Michigan Institute for Data Science, University of Michigan

**Abstract**

In the past few decades, mathematical models based on dynamical systems theory have provided new insight into diverse biological systems. In this review, we ask whether the recent success of machine learning techniques for large-scale biological data analysis provides a complementary or competing approach to more traditional modeling approaches. Recent applications of machine learning to the problem of learning biological dynamics in diverse systems range from neuroscience to animal behavior. We compare the underlying mechanisms and limitations of traditional dynamical models with those of machine learning models. We highlight the unique role that traditional modeling has played in providing predictive insight into biological systems, and propose several avenues for bridging traditional dynamical systems theory with large-scale analysis enabled by machine learning.

**Highlights**

- Machine learning, and particularly unsupervised learning, allows dynamics to readily be extracted from large biological datasets.
- Machine learning approaches can be combined with traditional mathematical modelling to yield better prediction of biological systems.
- Concepts originally developed for analyzing mathematical models may be used to interpret the properties of machine learning models.

**Introduction: The limitations of large mathematical models**

The most informative biological models unify observations across scales. For example, the celebrated Hodgkin-Huxley model of neural excitation comprises a set of differential equations describing the coupled dynamics of current, potential, and polarization within a single neuron's lipid bilayer (Hodgkin & Huxley 1952). Hodgkin and Huxley derived these equations from a first-principles analysis that combined the theory of electrical circuits with detailed anatomical study and measurements of squid giant axons. In their model, detailed experiments determined the form and parameters describing a single neuronal unit, and when many such units are combined together, their collective dynamics display the same emergent properties as large-scale neural ensembles (Figure 1). But while the Hodgkin-Huxley model was developed and validated in this specific experimental system, its legacy stems from its relevance to diverse neurobiological problems---ranging from the beating of cardiac myocytes, to alpha wave propagation in the human thalamus (Lynn et al 2019, Breakspear 2017). Similar "bottom-up" models recur throughout systems biology—for example, the Lotka-Volterra model of predator-prey competition may be scaled up to simulate entire multitrophic food webs (Lafferty et al 2015, Grilli et al 2017); and

the Goodwin model of negative feedback in an enzymatic network may be extended to describe large gene-regulatory networks (Forger 2017). In these cases, the structure of the model (equations) were determined directly from experimental observations.

However, the bottom-up approach has a significant limitation: while the individual terms in each dynamical equation are usually well-justified and mechanistic, precise measurements are often unavailable for every parameter---such as rate constants and saturation values describing interactions. Modellers instead resort to imposing strong assumptions on the parameters defining the model, such as by choosing parameters that cause the model to exhibit the correct large-scale collective behavior as observed in data. In rare circumstances, mathematical analysis can be used to show that the structure of the model yields properties that can be tested experimentally. In general, the accuracy of these parameter choices often remains untested.

Recent advances in machine learning challenge the classical “bottom up” approach to biological modelling. In the language of modern statistical learning, traditional modelling consists of model selection (defining an equation via mechanistic study), training (fitting rate constants and other parameters), and validation (prediction and confirmation of biologically-realistic dynamics). In machine learning, however, large-scale observations often serve as a starting point for analysis---all of the complexities and peculiarities of the observed data are visible to the model during training. Recently, “supervised” learning has generated significant interest among biologists due to remarkable advances in laboratory tasks requiring manual annotation of data---such as segmentation of cell images, or classification of mutations. From the perspective of understanding biological dynamics, however, we are particularly interested in *unsupervised* learning, which seeks to identify structure and patterns indirectly using unlabelled raw experimental data (e.g. motif discovery).

While many systems biology models blend aspects of bottom-up and top-down modelling, we believe that this dichotomy illustrates the potential power of modern statistical and machine learning to generate fundamentally new insights into systems biology. We argue that these techniques have the potential to work in tandem with traditional mathematical modeling techniques, revealing emergent simplicity in seemingly complex biological datasets.

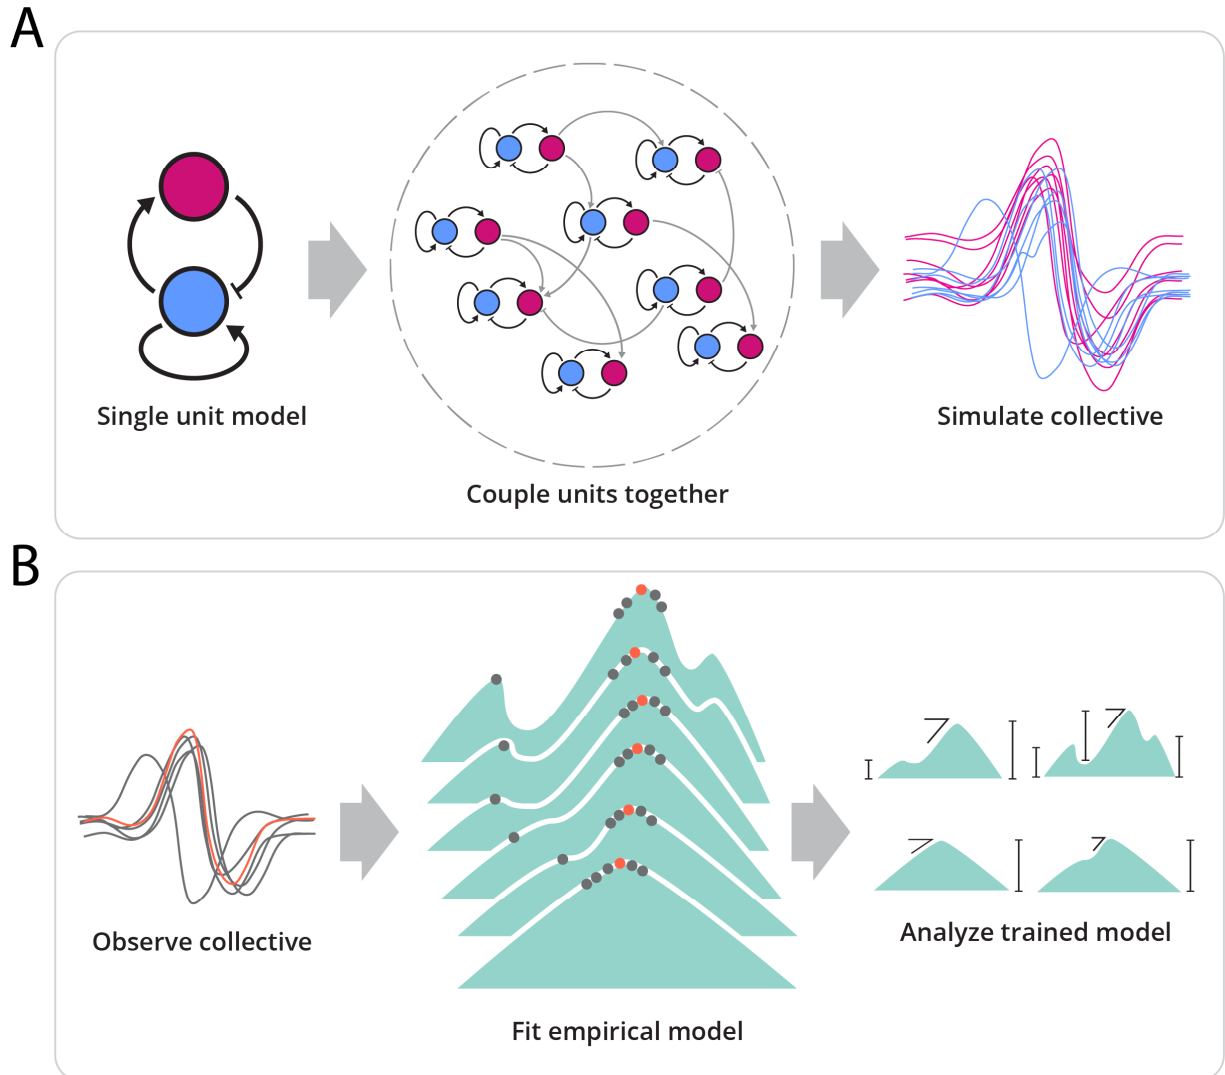


Figure 1. Bottom-up versus top-down modelling. (A) In bottom-up models, experimental data is used to derive a model of single units (e.g. chemicals, neurons, species). Many units are coupled together, and the resulting collective dynamics are analyzed. (B) In top-down models, experimental measurements of collective dynamics are measured directly, and then fit to an empirical model (here, a time-evolving probability distribution). The parameters and properties of the best-fit model (here, properties of the distribution) may then be analyzed and compared across timepoints or replicates.

### Notions of Time

Most classical models in systems biology are typically expressed in terms of differential equations of time. Given the initial state and input parameter values of a system, these equations tell us how the state of the system evolves over time. Machine learning models do not have an inherent concept of time, and so an auxiliary mapping between the model and time is required for proper comparison between these approaches. There are several ways to account for this:

- 1) One can run classical models until they reach an equilibrium state or endpoint, and then compare their input (initial conditions) and output (final state) to those achieved by machine learning approaches.
- 2) Some machine learning techniques take the full time-history of a system as an input for learning. These can be used to identify and classify different dynamical regimes within the time series, revealing different persistent behaviors (such as attractors).
- 3) Recurrent neural networks are a type of artificial neural network that explicitly include memory via an internal state, making them well-suited to modelling sequential data such as time series datasets.
- 4) The individual layers of a deep neural network can be considered as timepoints, in which case training the network via backpropagation of gradients across layers becomes equivalent to the classical adjoint method for optimizing the parameters of an ordinary differential equation---making it possible to define networks with a continuous layer index (Chen et al 2018). Furthering this analogy, Pontryagin's maximum principle (a concept from control theory) can be used to improve training of strongly constrained models (Zhang et al 2019, Weinan 2017)

Many linear mathematical models, or nonlinear mathematical models with small magnitude inputs, obey the superposition property: the effect of an input at one time is independent of the effect of an input at another time. Such dynamics can readily be incorporated into a machine learning model by introducing an appropriate regularization, which is a constraint that guides learning and deters overfitting. However, most mathematical models of biological systems do not have this property, and this distinction plays an important role in information processing in biological systems. For example, consider the processing of inputs by a single biological neuron. For inputs that balance excitation and inhibition, the most effective signal to generate firing is often a growing sinusoid with a period matched to the resonant period of the neuron (Forger 2017). However, for primarily inhibitory inputs, the most effective stimulus is a post-inhibitory rebound, or a sharp rise from inhibition. Moreover, a single neuron can also show different stable steady states (bistability) depending on its prior inputs. For example, it can exhibit quiescence or repetitive firing, and these states can persist until some other signal is strong enough to perturb the system state. Incorporating these properties into machine learning models requires explicit constraints to be imposed during training, while differential equation models naturally and concisely capture these properties.

### **Reconstructing governing coordinates from experimental data**

Various techniques based on unsupervised dimensionality reduction and manifold learning offer a promising application of machine learning to this problem. For example, researchers frequently record the simultaneous activity of thousands of neurons, under the implicit hypothesis that the dynamical variables necessary to model cognition may be reconstructed from these measurements alone (Marques et al 2020, Pandarinath et al 2018). These methods extract independent, low-dimensional governing coordinates from high-dimensional measurements using heuristics such as principal component analysis (PCA), which minimizes the covariance among extracted coordinates, or more sophisticated variants based on isolating timescales or recurring dynamical motifs (Hastie et al 2009). Unsupervised learning of dynamics has broad application to not only neural recordings (Pandarinath et al 2018) but also to microbial ecology (Tikhonov 2017), organismal locomotion (Brown 2018), and tracking animal behavior

(Brown et al 2018, Berman 2018, Mathis et al 2018). When large amounts of data are available, autoencoders (a type of neural network) are a particularly flexible technique for unsupervised learning (Figure 2). These networks seek to learn an identity function that maps the data onto itself, subject to constraints on how the data is stored and represented within the network. Autoencoders typically compress high-dimensional data into a low-dimensional “latent” space, and additional constraints (e.g. orthogonality, sparsity, continuity) may be introduced in order to influence the representation of the data within this latent space.

Similar techniques exist for classical biological models. Even when models contain hundreds or thousands of variables, the system may be accurately represented by a low-dimensional system whose variables are nonlinear combinations of the original system variables. An example of this is mathematical models of biological clocks, which contain many variables, but which can often be reduced to a subsystem (manifold) consisting of two variables (Forger & Kronauer 2002). Such low dimensional characterizations of the original system have variables and parameters that can be very complicated combinations of the original system variables. If one seeks biologically testable hypotheses about these models (e.g., if protein X increases in concentration, will protein Y increase as well?), the easiest path may be to simulate the full original system, rather than to try to determine this underlying smaller system.

Unsupervised learning can also be applied to the inverse problem of dimensionality *augmentation*, for cases in which the number of independent measurement channels is lower than that of a system’s underlying dynamics. For example, in an ecosystem of an interacting predator and prey, it may be possible to trap and measure only the predator---introducing the question of whether the dynamics of the prey can be inferred from measurements of the predator (Sugihara et al 2012). Similar “hidden variables” occur in molecular biology, in which time-resolved measurements of a handful of fluorescent reporters are used to infer the properties of a many-component gene regulatory network. In medical studies involving wearable sensors, low-dimensional time series measurements of skin conduction or accelerometer recordings are used to infer a subject’s high-dimensional behavioral or physiological state (Walch et al 2019, Huang et al 2020). Classical techniques for learning additional measurements—collectively known as “empirical dynamic modeling” or “state space reconstruction”—seek to use time series analysis to reconstruct the underlying state space or attractor of a system, allowing insights into laboratory or ecological systems for which observations are sparse (Clark et al 2020, Sugihara et al 2012). Recently, deep neural networks have been applied to this problem (Gilpin 2020), allowing the time series’ full history to be mined in order to produce robust and consistent proxies for unobserved dynamical variables.




Figure 2: A schematic of an autoencoder, a type of neural network with two halves: a “encoder,” which maps high dimensional input data to a low dimensional “latent” space, and a “decoder,” which maps the low-dimensional latent representation back to a high-dimensional reconstruction of the original input. The two halves are connected and trained together, producing a network that effectively maps the input data onto itself, subject to constraints on the dimensionality of the data’s representation within the latent space. After training, the decoder portion of the network may be set aside, and the encoder portion may be used to map additional input data into lower-dimensional latent variables.

### Merging machine learning with dynamical modelling

Recent studies have sought to combine large-scale analysis facilitated by unsupervised learning with the interpretability of low-dimensional dynamical systems models. For example, in a recent study of freely-behaving mice, an autoencoder was trained on 128 x 128 pixel videos of a mouse undergoing a range of free behaviors (Batty et al 2019). This autoencoder extracted ~10 latent variables that were sufficient to describe the majority of the mouse’s observed behavioral repertoire, and these latent variables were then used as inputs for a model of the mouse’s behavior as a stochastic dynamical system. Another recent study builds upon previously-developed mathematical models of human circadian timekeeping, and combines this analysis with machine learning to infer sleep stages based on recordings from wearable devices (Walch et al 2019).

Symbolic regression seeks to fit differential equations directly to observed data by drawing upon a dictionary of candidate basis functions (such as polynomials or trigonometric functions) (Schmidt et al 2009, Daniels et al 2015, Champion et al 2019, Udrescu & Tegmark 2020). This technique has been used to analyze problems ranging from interactions in biochemical networks (Casadiego et al 2017), to the outbreak dynamics of infectious diseases (Horrocks et al 2020) and the heat stimulus response of *C. elegans* (Daniels 2019). The choice and size of this function dictionary acts as a prior for the problem, with most methods seeking a balance between accuracy and parsimony. However, models found by symbolic regression are most revealing when the functions comprising the dictionary library have a mechanistic interpretation for the system under study. For example, in systems governed primarily by mechanical and geometric constraints (such as biomechanical systems), a best-fit differential equation model expressed in terms of trigonometric functions can reveal subtle mechanical symmetries and constraints (Schmidt et al 2009). Likewise, using a dictionary of Hill functions in a symbolic regression



model of a large gene regulatory network reveals the presence of specific governing interaction pathways and their rate laws (Daniels et al 2015). However, for datasets such as neural recordings, a readily-interpretable dictionary of functions describing the problem may not necessarily exist---limiting the applicability of symbolic regression.

For most biological systems, explicit but parsimonious enumeration of the dynamics in terms of known basis functions is likely impossible or computationally infeasible. These systems offer a promising platform for recent efforts to construct numerical approximations of time evolution operators, which computationally map an observation at one timepoint to its values at later timepoints (Lusch et al 2018, Pandarinath et al 2018). These approaches have proven particularly applicable to structural biology problems such as protein folding, where they allow informative coarse-grained models to be algorithmically extracted from detailed molecular dynamics simulations (Mardt et al 2018, Ramsundar et al 2019). These techniques are motivated by the general theoretical properties of nonlinear dynamical systems, which are difficult to analyze using traditional mathematical tools.

After fitting, numerical time evolution operators represent an empirical analogue of a differential equation conditioned on the training data. For example, a recent study used dynamic mode decomposition (a method originally developed for fluid dynamics) to identify transient resting states in a large dataset comprising fMRI readings of the human brain (Kunert et al 2019). Unlike PCA or other dimensionality reduction techniques, these methods associate dynamical motifs with scales in both space and time, providing more insight into spatiotemporal patterns than traditional decomposition techniques such as Fourier transforms. Other recent applications of this technique have been for behavior and gait quantification (Fujii et al 2019) and for optimization of reporter placements for cell state measurements (Hasnain et al 2019). While operator methods have not yet become widespread within the synthetic and systems biology communities, we anticipate that their strong theoretical underpinnings and efficient open-source implementations will facilitate their broad adoption.

### **Extrapolating predictions from limited data**

A lingering question when applying unsupervised learning to biological data is whether the appearance of low-dimensional effective coordinates (and a dynamical model defined on those coordinates) reveals any new information about the system itself. Traditional first-principle models posit a hypothesis in their formulation: the choice of terms in the governing equation represent a falsifiable theory of a system's underlying mechanism. But when low-dimensional coordinates are discovered by unsupervised learning, the interpolative nature of the resulting model represents a re-packaging of the existing data consistent with the constraints on the particular learning algorithm.

A key question in machine learning is whether the training data is sufficient to construct an accurate representation of the system for inputs that strongly differ from those encountered during training. In control theory, this is called the *identifiability* of a mathematical model. Only in select cases can this question be fully addressed. In a recent development, certain types of mathematical models can be shown to exhibit “completely determinable dynamics,” meaning that one recording of the system suffices to determine the behavior of a system for all other cases over a range of system states (Kim & Forger 2012,



Forger 2017). Thus, the behavior of the system over a finite time (a one-dimensional observation) determines the behavior over a range of system states (having the same dimension as the original system). However, not all systems have this property---for example, chaotic systems exhibit very different behavior for small changes in their state.

In dynamical systems theory, differential equation models are often analyzed by determining their equilibrium points, and then performing stability analysis of the model at these fixed points. In machine learning, a similar technique is saliency mapping, which calculates the approximate derivative of a trained model's outputs with respect to its inputs (Hastie et al 2009). This process identifies the input features that most strongly determine the model's predictions. For example, in a recent study a neural network was trained to predict the crawling trajectories of developing muscle cells (Kimmel et al 2019). Saliency maps were calculated for the trained model with respect to the input time series, and the maps detected unusual bouts of locomotion that differentiated cell types---thereby identifying qualitative traits that microscopists could search for when observing cells. Another common technique for analyzing differential equation models is sensitivity analysis, which seeks to determine how a model's outputs change as the parameters describing individual terms in the model change. For a machine learning model, a comparable technique is ablation, in which single components are removed from a trained model (such as architectural features of a neural network) in order to identify which terms in the model most strongly influence its accuracy (Goodfellow et al 2016).

## Conclusion and outlook

We have shown diverse examples demonstrating the ability of machine learning models to distill insights into large-scale biological dynamics. Does the ability of unsupervised learning to identify low-dimensional coordinates in diverse biological systems suggest deeper biological principles?

A counterintuitive result of the ever-larger-scale deployment of artificial neural networks is the observation that many aspects of large models are surprisingly consistent across different implementations and applications (Bahri et al 2020). Despite having thousands, or even millions, of trainable parameters, artificial neural networks often converge to similar solutions across replicates, to the degree that their training can sometimes be approximated by low-order ordinary differential equations (Jacot et al 2018, Bordelon et al 2020). Such results mirror the phenomenology of large-scale biological regulatory networks, which, despite comprising many interacting genes and proteins, often have a small set of effective governing parameters---such as the kinetic constants of a few rate-limiting reactions---that jointly determine the system's large-scale dynamics (Gutenkunst et al 2007). Similar results are observed in large ecological communities, for which isolated study of individual species---and thus validated models of individual interactions---is difficult to achieve (Gross et al 2009). In neuroscience, many theoretical and experimental studies demonstrate cases in which thousands of interacting neurons collectively produce low-dimensional dynamics (Lynn et al 2019, Forger 2017, Bick et al 2020). For example, by applying an unsupervised learning algorithm based on topological filtration, the spiking activity of thousands of neurons in a mouse thalamus can be robustly mapped onto a low-dimensional "attractor manifold" encoding the angular bearing of the animal's head (Chaudhuri et al 2019). Intriguingly, direct comparison of neural activity in the visual system to activation patterns in trained

convolutional neural networks (the current state-of-the-art technique for image analysis) suggest that artificial and biological neural networks may even share common hierarchical representation and organization of image information (Jaegle et al 2019, Bashivan et al 2019).

Taken together, these results introduce the question of whether universal mathematical constraints determine certain aspects of large biological systems---that large, overparameterized systems of weakly-interacting units spontaneously collapse onto a low-dimensional manifold. Has evolution driven complex biological networks towards these emergent motifs, and do they confer adaptive benefits from the perspective of control or stability---such as stabilizing an ecosystem against an invasive predator, or suppressing unwanted fluctuations in a large genetic circuit? We hope that further development of models at the intersection of machine learning and dynamical systems theory will provide unified insight into this question.

**Acknowledgements:** We thank the NSF-Simons Center for Mathematical and Statistical Analysis of Biology at Harvard University, NSF Grant No. DMS-1764269, and the Harvard FAS Quantitative Biology Initiative.

## References

Udrescu, Silviu-Marian, and Max Tegmark. "AI Feynman: a physics-inspired method for symbolic regression." *Science Advances* (2020).

Forger, Daniel B., and Richard E. Kronauer. "Reconciling mathematical models of biological clocks by averaging on approximate manifolds." *SIAM Journal on Applied Mathematics* 62.4 (2002): 1281-1296.

Zhang, Dinghui, et al. "You only propagate once: Accelerating adversarial training via maximal principle." *Advances in Neural Information Processing Systems*. 2019.

Lusch, Bethany, J. Nathan Kutz, and Steven L. Brunton. "Deep learning for universal linear embeddings of nonlinear dynamics." *Nature communications* 9.1 (2018): 1-10.

Weinan, E. "A proposal on machine learning via dynamical systems." *Communications in Mathematics and Statistics* 5.1 (2017): 1-11.

**Champion, Kathleen, et al. "Data-driven discovery of coordinates and governing equations." *Proceedings of the National Academy of Sciences* 116.45 (2019): 22445-22451.**

**++Demonstrates a systematic machine learning technique for extracting governing coordinates from spatiotemporal datasets, and then fitting the resulting coordinates to differential equations using sparse symbolic regression.**

Batty, Eleanor, et al. "BehaveNet: nonlinear embedding and Bayesian neural decoding of behavioral videos." *Advances in Neural Information Processing Systems*. 2019.

Gutenkunst, Ryan N., et al. "Universally sloppy parameter sensitivities in systems biology models." *PLoS computational biology* 3.10 (2007).

Hastie, Trevor, Robert Tibshirani, and Jerome Friedman. *The elements of statistical learning: data mining, inference, and prediction*. Springer Science & Business Media, 2009.

Jacot, Arthur, Franck Gabriel, and Clément Hongler. "Neural tangent kernel: Convergence and generalization in neural networks." *Advances in Neural Information Processing Systems*. 2018.

Bordelon, Blake, Abdulkadir Canatar, and Cengiz Pehlevan. "Spectrum Dependent Learning Curves in Kernel Regression and Wide Neural Networks." *International Conference on Machine Learning* (2020).

Ramsundar, Bharath, Peter Eastman, Patrick Walters, and Vijay Pande. *Deep learning for the life sciences: applying deep learning to genomics, microscopy, drug discovery, and more*. " O'Reilly Media, Inc.", 2019.

Gilpin, William. "Cellular automata as convolutional neural networks." *Physical Review E* 100.3 (2019): 032402.

Bick, Christian, et al. "Understanding the dynamics of biological and neural oscillator networks through mean-field reductions: A review." *The Journal of Mathematical Neuroscience* (2020).

Lynn, Christopher W., and Danielle S. Bassett. "The physics of brain network structure, function and control." *Nature Reviews Physics* 1.5 (2019): 318.

**Chen, Tian Qi, et al. "Neural ordinary differential equations." *Advances in Neural Information Processing Systems*. 2018.**

**+ Identifies an equivalence between backpropagation of errors across layers in deep neural networks, and the adjoint method used for optimization of parameters in ODEs and PDEs.**

Lafferty, Kevin D., et al. "A general consumer-resource population model." *Science* 349.6250 (2015): 854-857.

Kim, Jae K., and Daniel B. Forger. "On the existence and uniqueness of biological clock models matching experimental data." *SIAM Journal on Applied Mathematics* 72.6 (2012): 1842-1855.

**Forger, Daniel B. *Biological clocks, rhythms, and oscillations: the theory of biological timekeeping*. MIT Press, 2017.**

**+An in-depth study of modeling and analysis of biological rhythms.**

Gross, Thilo, et al. "Generalized models reveal stabilizing factors in food webs." *Science* 325.5941 (2009): 747-750.

Breakspear, Michael. "Dynamic models of large-scale brain activity." *Nature neuroscience* 20.3 (2017): 340.

Mezić, Igor. "Analysis of fluid flows via spectral properties of the Koopman operator." *Annual Review of Fluid Mechanics* 45 (2013): 357-378.

Brown, André EX, and Benjamin De Bivort. "Ethology as a physical science." *Nature Physics* 14.7 (2018): 653-657.

Pathak, Jaideep, et al. "Model-free prediction of large spatiotemporally chaotic systems from data: A reservoir computing approach." *Physical review letters* 120.2 (2018): 024102.

**Bashivan, Pouya, Kohitij Kar, and James J. DiCarlo. "Neural population control via deep image synthesis." *Science* 364.6439 (2019): eaav9436.**

**++The authors design an artificial neural network to mimic the visual system, and then use it to generate images that selectively activate different neuron populations in the macaque visual cortex. Using various metrics, the authors found that the artificial neural network reproduced several key aspects of the animals' neural activity patterns.**

Jaegle, Andrew, et al. "Population response magnitude variation in inferotemporal cortex predicts image memorability." *eLife* 8 (2019).

Mathis, Alexander, et al. "DeepLabCut: markerless pose estimation of user-defined body parts with deep learning." *Nature neuroscience* 21.9 (2018): 1281.

Fujii, Keisuke, et al. "Data-driven spectral analysis for coordinative structures in periodic human locomotion." *Scientific reports* 9.1 (2019): 1-14.

Goodfellow, Ian, Yoshua Bengio, and Aaron Courville. *Deep learning*. MIT press, 2016.

Berman, Gordon J. "Measuring behavior across scales." *BMC biology* 16.1 (2018): 23.

Mardt, Andreas, et al. "VAMPnets for deep learning of molecular kinetics." *Nature Communications* 9.1 (2018): 1-11.

Gilpin, William. "Deep reconstruction of strange attractors from time series." *arXiv preprint arXiv:2002.05909* (2020).

Grilli, Jacopo, et al. "Higher-order interactions stabilize dynamics in competitive network models." *Nature* 548.7666 (2017): 210-213.

Clark, T. J., and Angela D. Luis. "Nonlinear population dynamics are ubiquitous in animals." *Nature Ecology & Evolution* 4.1 (2020): 75-81.

- Pandarinath, Chethan, et al. "Inferring single-trial neural population dynamics using sequential auto-encoders." *Nature methods* 15.10 (2018): 805-815.
- Lee, Jaehoon, et al. "Wide neural networks of any depth evolve as linear models under gradient descent." *Advances in neural information processing systems*. 2019.
- Horrocks, Jonathan, and Chris T. Bauch. "Algorithmic discovery of dynamic models from infectious disease data." *Scientific Reports* 10.1 (2020): 1-18.
- Hodgkin, Alan L., and Andrew F. Huxley. "A quantitative description of membrane current and its application to conduction and excitation in nerve." *The Journal of physiology* 117.4 (1952): 500-544.
- Daniels, Bryan C., and Ilya Nemenman. "Automated adaptive inference of phenomenological dynamical models." *Nature communications* 6.1 (2015): 1-8.
- Schmidt, Michael, and Hod Lipson. "Distilling free-form natural laws from experimental data." *Science* 324.5923 (2009): 81-85.
- Kunert-Graf, James Michael, et al. "Extracting reproducible time-resolved resting state networks using dynamic mode decomposition." *Frontiers in computational neuroscience* 13 (2019): 75.
- Casadio, Jose, et al. "Model-free inference of direct network interactions from nonlinear collective dynamics." *Nature communications* 8.1 (2017): 1-10.
- Daniels, Bryan C., William S. Ryu, and Ilya Nemenman. "Automated, predictive, and interpretable inference of *Caenorhabditis elegans* escape dynamics." *Proceedings of the National Academy of Sciences* 116.15 (2019): 7226-7231.
- Sugihara, George, et al. "Detecting causality in complex ecosystems." *Science* 338.6106 (2012): 496-500.
- Bahri, Yasaman, et al. "Statistical mechanics of deep learning." *Annual Review of Condensed Matter Physics* (2020).
- Marques, João C., et al. "Internal state dynamics shape brainwide activity and foraging behaviour." *Nature* 577.7789 (2020): 239-243.
- Hasnain, Aqib, Nibodh Boddupalli, and Enoch Yeung. "Optimal reporter placement in sparsely measured genetic networks using the Koopman operator." *arXiv preprint arXiv:1906.00944* (2019).
- Walch, Olivia J., Yitong Huang, Daniel B. Forger, and Cathy Goldstein. "Sleep stage prediction with raw acceleration and photoplethysmography heart rate data derived from a consumer wearable device." *Sleep* 42 (2019): zsz180.

Huang, Yitong et al. "Predicting circadian phase across populations: a comparison of mathematical models and wearable devices." Under review

**Declaration of interests**

☒ The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

☐ The authors declare the following financial interests/personal relationships which may be considered as potential competing interests: